

# Regularisation notes.

Robert K.

December 27, 2016

## Abstract

This is just a document with notes regarding Regularisation in finding  $\Theta$  values to prevent *overfitting*.

## 1 Underfit vs Overfit

*Underfit* - (a.k.a "high bias") is the prediction (learned hypothesis function) which does not fit the training or new input data very well. Too simplified hypothesis function.

*Overfit* - (a.k.a "high variance") is the prediction (learned hypothesis function) which fits well the training data, but fails to generalise to new input data. It can occur when we have too many *features* and small size of training set.

How to address overfitting problem:

1. Reduce the number of features, which can be done in two ways:
  - Manually select which features to keep
  - Use Model Selection Algorithm to automatically select features to keep
2. Use **regularisation** where we keep all features, but with different priority

## 2 Regularisation

**Regularisation** - is the technique of reducing overfitting by keeping all features but reducing the values of  $\theta_j$ .

It works well when we have a lot of features and each of them contributes a bit to predicting  $y$ .

*Properties of regularisation:*

- if  $\theta_j$  is small then more likely we will not overfit. We want to keep  $\theta_j$  small.

## 2.1 Regularisation of linear regression

### 2.1.1 Cost function and gradient

To use Regularisation we slightly modify cost function  $J(\theta)$

$$\begin{aligned} J(\theta) &= \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right] = \\ &= \left[ \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \right] + \left[ \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2 \right] \end{aligned} \quad (1)$$

If the  $\theta_j$  is small, then the cost function  $J(\theta)$  will be small too.  
 $\lambda$  - is the **regularisation parameter** It determines how much the costs of our theta parameters are inflated.

*Properties of new cost function:*

- cost will be small if  $\theta_j$  is small.
- cost will be small for the big **regularisation parameter**  $\lambda$  if  $\theta_j$  will be super small
- If  $\lambda$  is chosen to be too large, then  $\theta_1 \dots \theta_n$  will be near zero to make cost small. It may cause underfitting.
- if  $\lambda = 0$  or is too small then it is similar to case of not using regularisation at all - so we increase a chance of overfitting.

Gradient of regularised cost function  $\frac{\delta}{\delta \theta_j} J(\theta)$

$$\frac{\delta}{\delta \theta_j} J(\theta) = \frac{\delta}{\delta \theta_j} \left[ \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \right] + \frac{\delta}{\delta \theta_j} \left[ \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2 \right] \quad (2)$$

where:

$$h_{\theta}(x^{(i)}) = \theta^T x^{(i)} = \sum_{j=0}^n \theta_j x_j^{(i)} \text{ with } x_0^{(i)} = 1$$

so gradient of cost function for linear regression:

$$\frac{\delta}{\delta \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \text{ for } j = 0 \quad (3)$$

$$\frac{\delta}{\delta \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \text{ for } j \geq 1 \quad (4)$$

The **vectorised form of cost function**

$$J(\Theta) = J \left( \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix} \right) = \frac{1}{2m} (X\Theta - y)^T (X\Theta - y) + \frac{\lambda}{2m} \begin{bmatrix} 0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}^T \begin{bmatrix} 0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix} \quad (5)$$

The **vectorised form of cost function derivative**

$$\nabla J(\Theta) = \begin{bmatrix} \frac{\delta}{\delta \Theta_0} J(\Theta) \\ \frac{\delta}{\delta \Theta_1} J(\Theta) \\ \vdots \\ \frac{\delta}{\delta \Theta_n} J(\Theta) \end{bmatrix} = \frac{1}{m} [X^T (X\Theta - y)] + \frac{\lambda}{m} \begin{bmatrix} 0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix} \quad (6)$$

### 2.1.2 Regularised normal equations

Regularisation using the non-iterative normal equation.

$$\Theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix} = \theta = \left( X^T X + \lambda \begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{bmatrix} \right)^{-1} X^T y \quad (7)$$

## 2.2 Regularisation of logistic regression

### 2.2.1 Cost function and gradient

Regularised **cost function** has a form

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_\theta(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2 \quad (8)$$

where:

$$y \in \{0, 1\}$$

$$h_\theta(x^{(i)}) = \frac{1}{1 + e^{-\theta^T x^{(i)}}}$$

**Derivative** of the cost function for logistic regression looks like that:

$$\frac{\delta}{\delta \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)} \text{ for } j = 0 \quad (9)$$

$$\frac{\delta}{\delta \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \text{ for } j \geq 0 \quad (10)$$

where:

$$y \in \{0, 1\}$$

$$h_{\theta}(x^{(i)}) = \frac{1}{1 + e^{-\theta^T x^{(i)}}}$$

The **vectorised form of cost function**

$$J(\Theta) = J \left( \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix} \right) = -\frac{1}{m} [y^T \log(g(X\Theta)) + (1 - y)^T \log(1 - g(X\Theta))] + \frac{\lambda}{2m} \begin{bmatrix} 0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}^T \begin{bmatrix} 0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix} \quad (11)$$

where:

$$g(X\Theta) = \begin{bmatrix} g(\Theta^T x^{(1)}) \\ g(\Theta^T x^{(2)}) \\ \vdots \\ g(\Theta^T x^{(m)}) \end{bmatrix}$$

$\log(g(X\Theta))$  - is a column vector of size  $m$ .  $\log$  function is done on every element of  $g(X\Theta)$  column vector.

The **vectorised form of cost function derivative**

$$\nabla J(\Theta) = \begin{bmatrix} \frac{\delta}{\delta \Theta_0} J(\Theta) \\ \frac{\delta}{\delta \Theta_1} J(\Theta) \\ \vdots \\ \frac{\delta}{\delta \Theta_n} J(\Theta) \end{bmatrix} = \frac{1}{m} [X^T (g(X\Theta) - y)] + \frac{\lambda}{m} \begin{bmatrix} 0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix} \quad (12)$$

where:

$$g(X\Theta) = \begin{bmatrix} g(\Theta^T x^{(1)}) \\ g(\Theta^T x^{(2)}) \\ \vdots \\ g(\Theta^T x^{(m)}) \end{bmatrix}$$